

Lecture 01

Bits, Bytes and Integers – part 1

Euhyun Moon, Ph.D.
Machine Learning Systems (MLSys) Lab
Computer Science and Engineering
Sogang University



Slides adapted from Randy Bryant and Dave O'Hallaron: Introduction to Computer Systems, CMU

Boolean Algebra

- Developed by George Boole in 19th Century
 - Algebraic representation of logic (True \rightarrow 1, False \rightarrow 0)
 - AND: $A \& B = 1$ when both A is 1 and B is 1
 - OR: $A | B = 1$ when either A is 1 or B is 1
 - XOR: $A \wedge B = 1$ when either A is 1 or B is 1, but not both
 - NOT: $\sim A = 1$ when A is 0 and vice-versa
 - DeMorgan's Law:
$$\sim (A | B) = \sim A \& \sim B$$
$$\sim (A \& B) = \sim A | \sim B$$

AND			OR			XOR			NOT	
$\&$	0	1	$ $	0	1	\wedge	0	1	\sim	
0	0	0	0	0	1	0	0	1	0	1
1	0	1	1	1	1	1	1	0	1	0

General Boolean Algebras

- Operate on bit vectors
 - Operations applied bitwise
 - All of the properties of Boolean algebra apply

01101001	01101001	01101001	
& 01010101	01010101	^ 01010101	~ 01010101

- Examples of useful operations:

$$x \wedge x = 0$$

	01010101
^	01010101
<hr/>	
	00000000

$$x | 1 = 1, \quad x | 0 = x$$

	01010101
	1 1110000
<hr/>	
	11110101

Bit-Level Operations in C

- **& (AND), | (OR), ^ (XOR), ~ (NOT)**

- View arguments as bit vectors, apply operations bitwise
- Apply to any “integral” data type
 - long, int, short, char, unsigned

- **Examples with char a, b, c;**

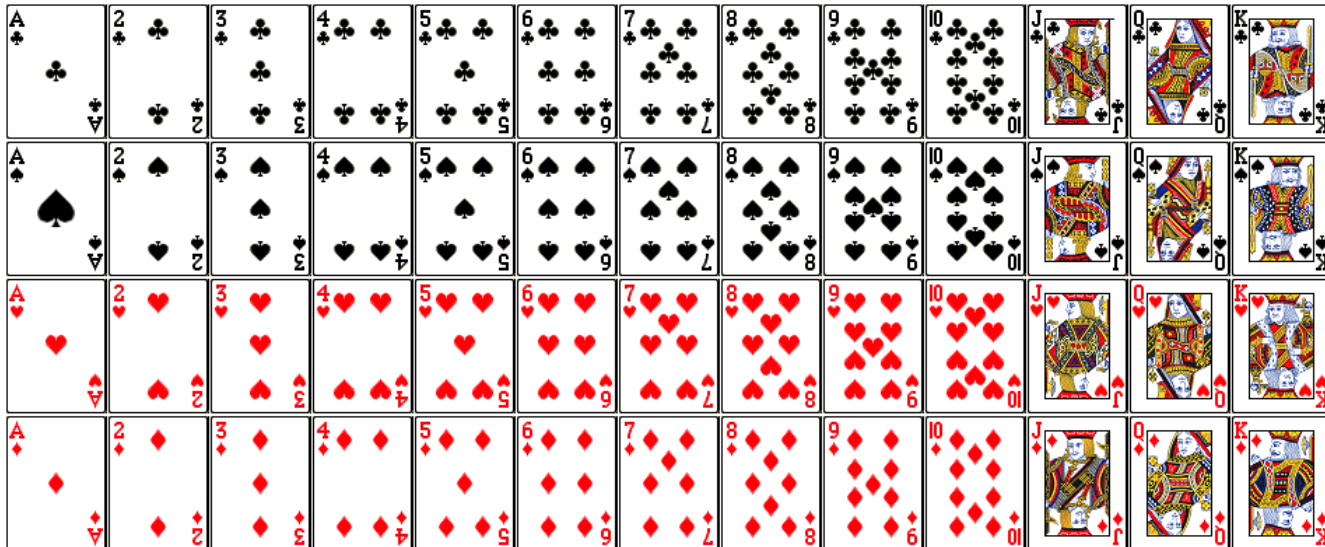
- ```
a = (char) 0x41; // 0x41 → 0b 0100 0001
b = ~a; // → 0x
```
- ```
a = (char) 0x69;    // 0x69 → 0b 0110 1001
b = (char) 0x55;    // 0x55 → 0b 0101 0101
c = a & b;          //                → 0x
```
- ```
a = (char) 0x41; // 0x41 → 0b 0100 0001
b = a; // 0b 0100 0001
c = a ^ b; // → 0x
```

# Contrast: Logic Operations in C

- Logical operators in C: **&& (AND)**, **|| (OR)**, **! (NOT)**
  - **0** is False, **anything nonzero** is True
  - **Always** return 0 or 1
  - **Early termination** (a.k.a. short-circuit evaluation) of **&&**, **||**
- Examples (`char` data type)
  - `!0x41`  $\rightarrow$  `0x00`
  - `!0x00`  $\rightarrow$  `0x01`
  - `!!0x41`  $\rightarrow$  `0x01`
  - `p && *p`
    - If `p` is the **null pointer** (`0x0`), then `p` is never dereferenced!
  - `0xCC && 0x33`  $\rightarrow$  `0x01`
  - `0x00 || 0x33`  $\rightarrow$  `0x01`

# But before we get to integers....

- Encode a standard deck of playing cards
  - 52 cards in 4 suits
    - How do we encode suits, face cards?
- What operations do we want to make easy to implement?
  - Which is the higher value card?
  - Are they the same suit?



# Two possible representations

1) 1 bit per card (52): bit corresponding to card set to 1



low-order 52 bits of 64-bit word

- “One-hot” encoding (similar to set notation)
- Drawbacks:
  - Hard to compare values and suits
  - Large number of bits required

2) 1 bit per suit (4), 1 bit per number (13): 2 bits set



4 suits      13 numbers

- Pair of one-hot encoded values
- Easier to compare suits and values, but still lots of bits used

# Two better representations

## 3) Binary encoding of all 52 cards – only 6 bits needed

- $2^6 = 64 \geq 52$



low-order 6 bits of a byte

- Fits in one byte (smaller than one-hot encodings)
- How can we make value and suit comparisons easier?





## 4) Separate binary encodings of suit (2 bits) and value (4 bits)



suit value

- Also fits in one byte, and easy to do comparisons

| K    | Q    | J    | ... | 3    | 2    | A    |
|------|------|------|-----|------|------|------|
| 1101 | 1100 | 1011 | ... | 0011 | 0010 | 0001 |

|                                                                                       |    |
|---------------------------------------------------------------------------------------|----|
|   | 00 |
|  | 01 |
|  | 10 |
|  | 11 |



# Compare Card Suits

**mask:** a bit vector designed to achieve a desired behavior when used with a bitwise operator on another bit vector  $v$ . Here we turn all *but* the bits of interest in  $v$  to 0.

```
char hand[5]; // represents a 5-card hand
char card1, card2; // two cards to compare
card1 = hand[0];
card2 = hand[1];
...
if (sameSuitP(card1, card2)) { ... }
```

```
#define SUIT_MASK 0x30
```

```
int sameSuitP(char card1, char card2) {
 return (!((card1 & SUIT_MASK) ^ (card2 & SUIT_MASK)));
 //return (card1 & SUIT_MASK) == (card2 & SUIT_MASK);
}
```

returns **int**

SUIT\_MASK = 0x30 =

|   |   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|---|
| 0 | 0 | 1 | 1 | 0 | 0 | 0 | 0 |
|---|---|---|---|---|---|---|---|

suit      value

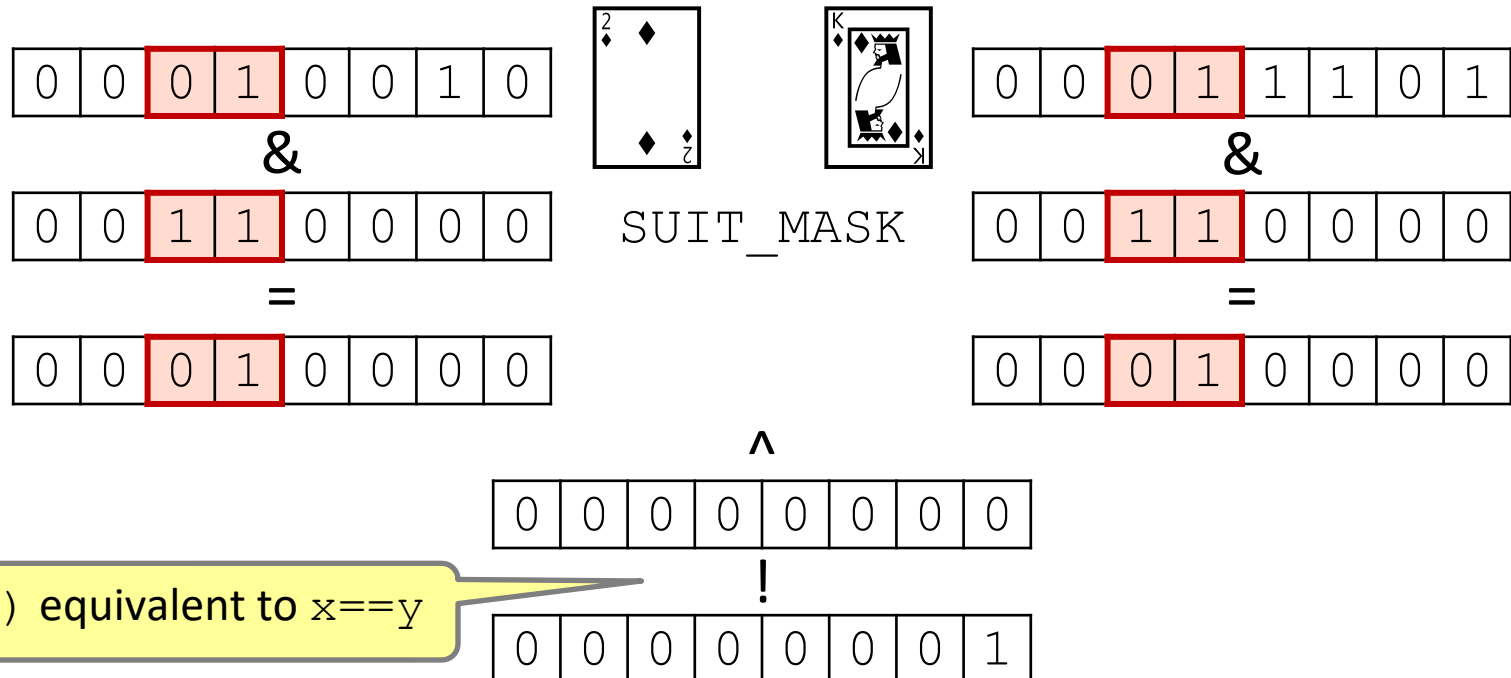
equivalent

# Compare Card Suits

**mask:** a bit vector designed to achieve a desired behavior when used with a bitwise operator on another bit vector  $v$ . Here we turn all *but* the bits of interest in  $v$  to 0.

```
#define SUIT_MASK 0x30
```

```
int sameSuitP(char card1, char card2) {
 return (!((card1 & SUIT_MASK) ^ (card2 & SUIT_MASK)));
 //return (card1 & SUIT_MASK) == (card2 & SUIT_MASK);
}
```



# Compare Card Values

**mask:** a bit vector designed to achieve a desired behavior when used with a bitwise operator on another bit vector *v*.

```
char hand[5]; // represents a 5-card hand
char card1, card2; // two cards to compare
card1 = hand[0];
card2 = hand[1];
...
if (greaterValue(card1, card2)) { ... }
```

```
#define VALUE_MASK 0x0F

int greaterValue(char card1, char card2) {
 return ((unsigned int)(card1 & VALUE_MASK) >
 (unsigned int)(card2 & VALUE_MASK));
}
```

VALUE\_MASK = 0x0F = 

|   |   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 1 | 1 | 1 | 1 |
|---|---|---|---|---|---|---|---|

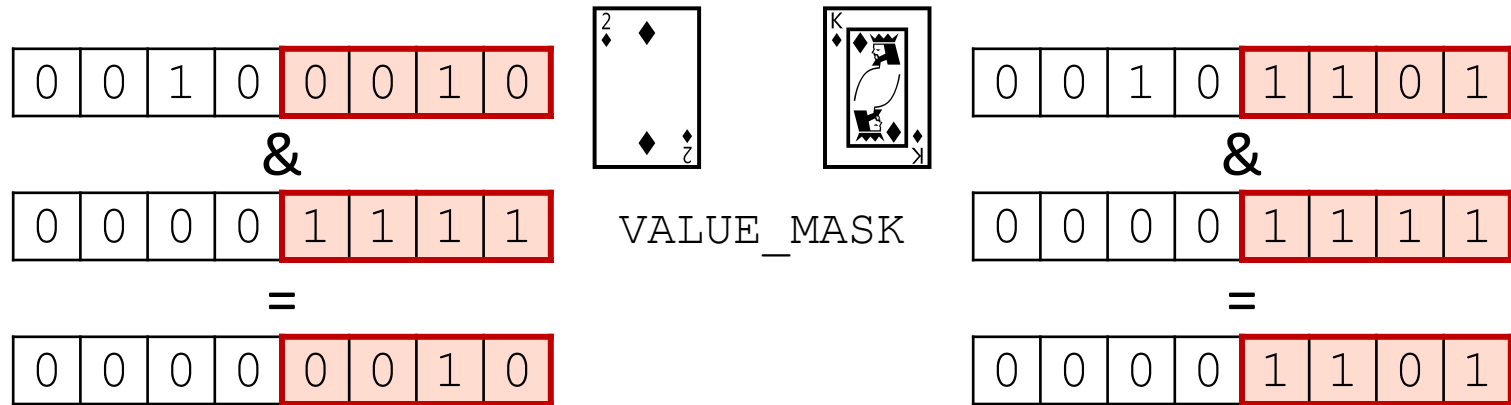
                                
suit                  value

# Compare Card Values

**mask:** a bit vector designed to achieve a desired behavior when used with a bitwise operator on another bit vector *v*.

```
#define VALUE_MASK 0x0F

int greaterValue(char card1, char card2) {
 return ((unsigned int)(card1 & VALUE_MASK) >
 (unsigned int)(card2 & VALUE_MASK));
}
```



$2_{10} > 13_{10}$

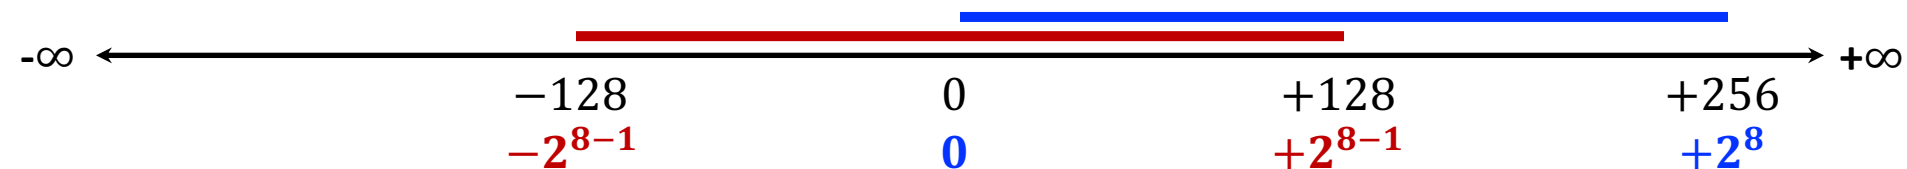
0 (false)

# Integers

- **Binary representation of integers**
  - **Unsigned and signed**
- Shifting and arithmetic operations
- In C: Signed, Unsigned and Casting
- Consequences of finite width representations
  - Overflow, sign extension

# Encoding Integers

- The hardware (and C) supports two flavors of integers
  - *unsigned* – only the non-negatives
  - *signed* – both negatives and non-negatives
- Cannot represent all integers with  $w$  bits
  - Only  $2^w$  distinct bit patterns
  - Unsigned values:  $0 \dots 2^w - 1$
  - Signed values:  $-2^{w-1} \dots 2^{w-1} - 1$
- **Example:** 8-bit integers (*e.g.*, `char`)



# 1) Unsigned Integers

- Unsigned values follow the standard base 2 system
  - $b_7b_6b_5b_4b_3b_2b_1b_0 = b_72^7 + b_62^6 + \dots + b_12^1 + b_02^0$
- Add and subtract using the normal “carry” and “borrow” rules, just in binary

|                                                       |                                                                          |
|-------------------------------------------------------|--------------------------------------------------------------------------|
| $\begin{array}{r} 63 \\ + 8 \\ \hline 71 \end{array}$ | $\begin{array}{r} 00111111 \\ + 00001000 \\ \hline 01000111 \end{array}$ |
|-------------------------------------------------------|--------------------------------------------------------------------------|

- Useful formula:  $2^{N-1} + 2^{N-2} + \dots + 2 + 1 = 2^N - 1$ 
  - *i.e.*, N ones in a row =  $2^N - 1$
- How would you make *signed* integers?

## 2) Sign and Magnitude

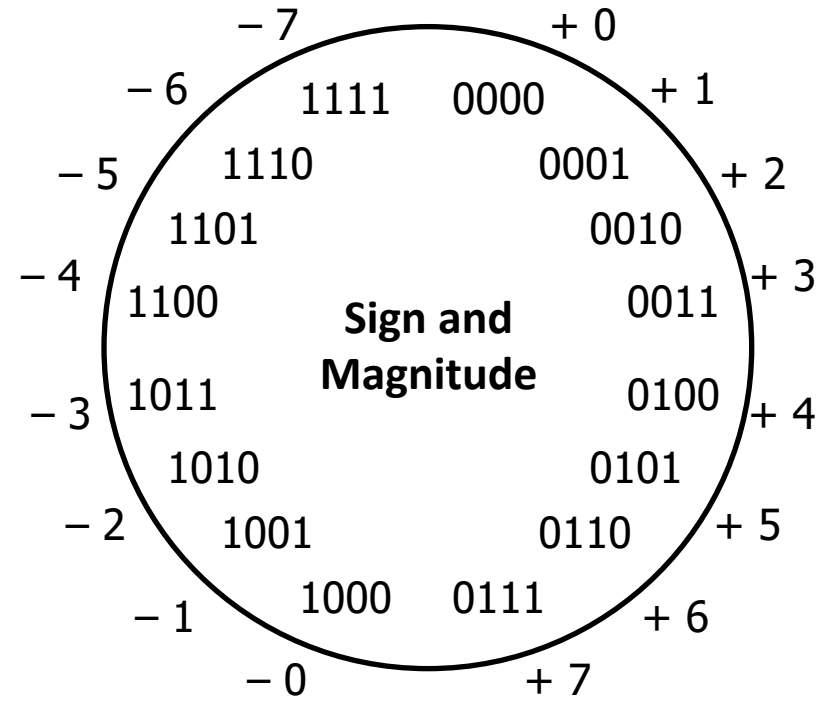
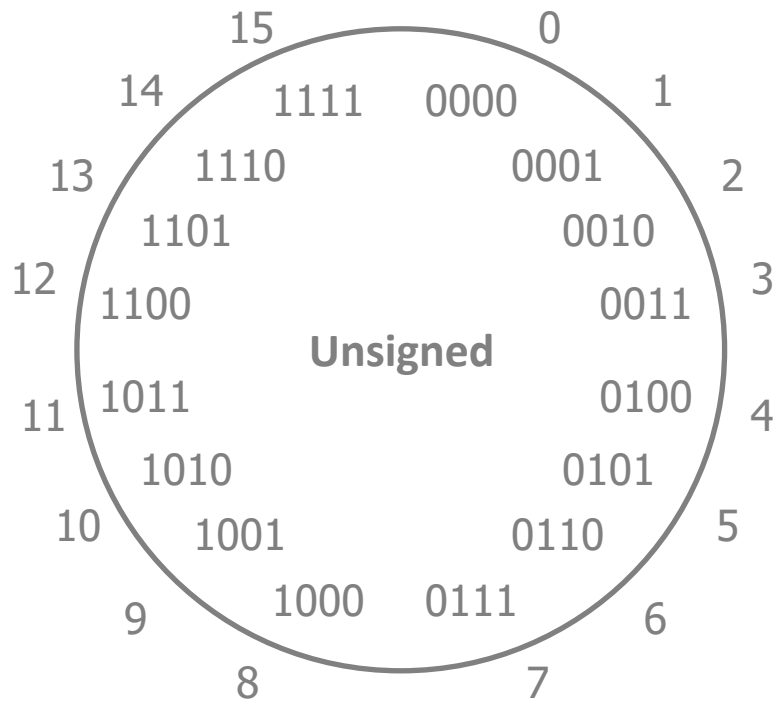
Most Significant Bit

- Designate the high-order bit (MSB) as the “sign bit”
  - $\text{sign}=0$ : positive numbers;  $\text{sign}=1$ : negative numbers
- Benefits:
  - Using MSB as sign bit matches positive numbers with unsigned
  - All zeros encoding is still = 0
- Examples (8 bits):
  - $0x00 = 00000000_2$  is non-negative, because the sign bit is 0
  - $0x7F = 01111111_2$  is non-negative ( $+127_{10}$ )
  - $0x85 = 10000101_2$  is negative ( $-5_{10}$ )
  - $0x80 = 10000000_2$  is negative... ... zero???



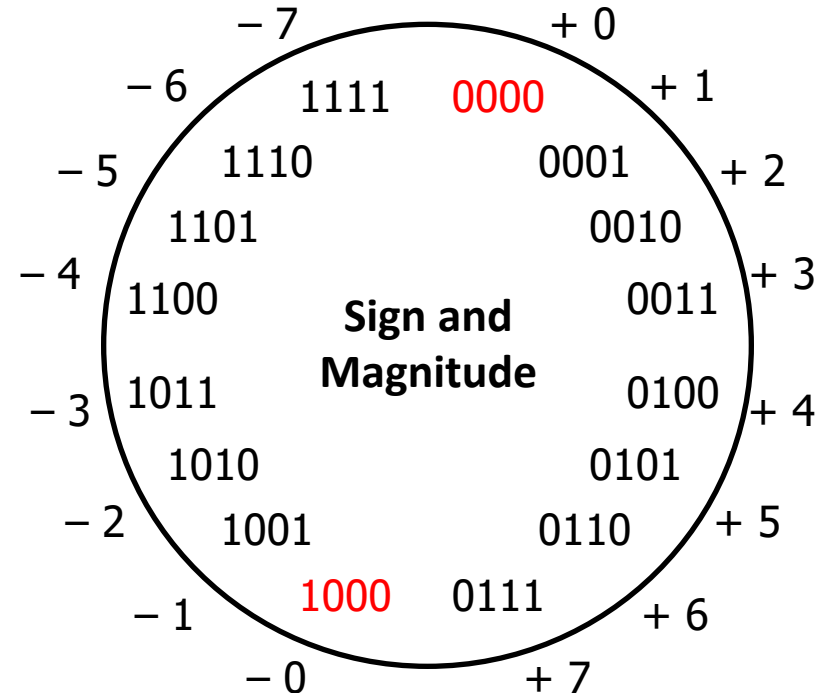
## 2) Sign and Magnitude

- MSB is the sign bit, rest of the bits are magnitude
- Drawbacks?



## 2) Sign and Magnitude

- MSB is the sign bit, rest of the bits are magnitude
- Drawbacks:
  - Two representations of 0 (bad for checking equality)



## 2) Sign and Magnitude

- MSB is the sign bit, rest of the bits are magnitude
- Drawbacks:
  - Two representations of 0 (bad for checking equality)
  - **Arithmetic is cumbersome**
    - Example:  $4 - 3 \neq 4 + (-3)$

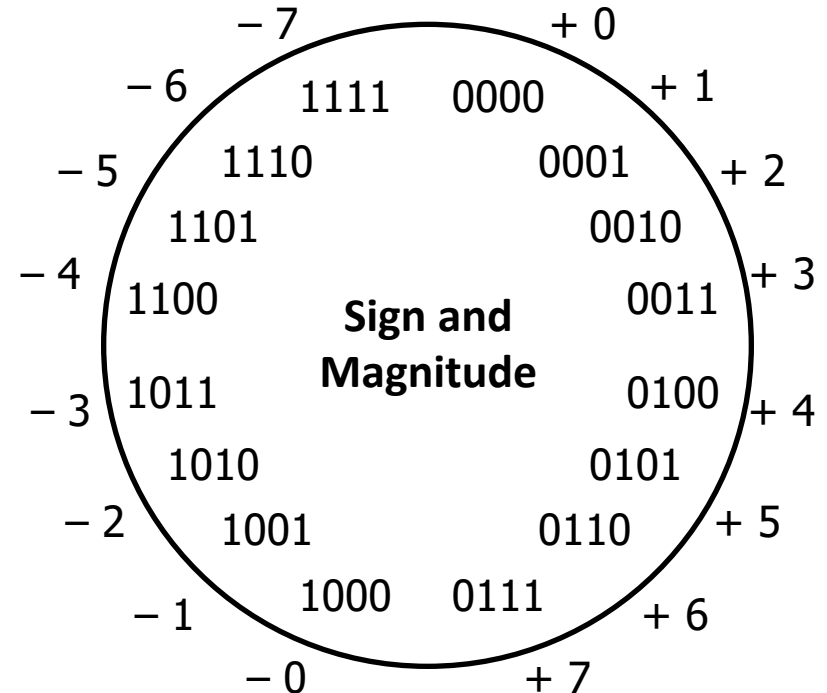
|       |        |
|-------|--------|
| 4     | 0100   |
| - 3   | - 0011 |
| <hr/> |        |
| 1     | 0001   |



|       |        |
|-------|--------|
| 4     | 0100   |
| + -3  | + 1011 |
| <hr/> |        |
| -7    | 1111   |

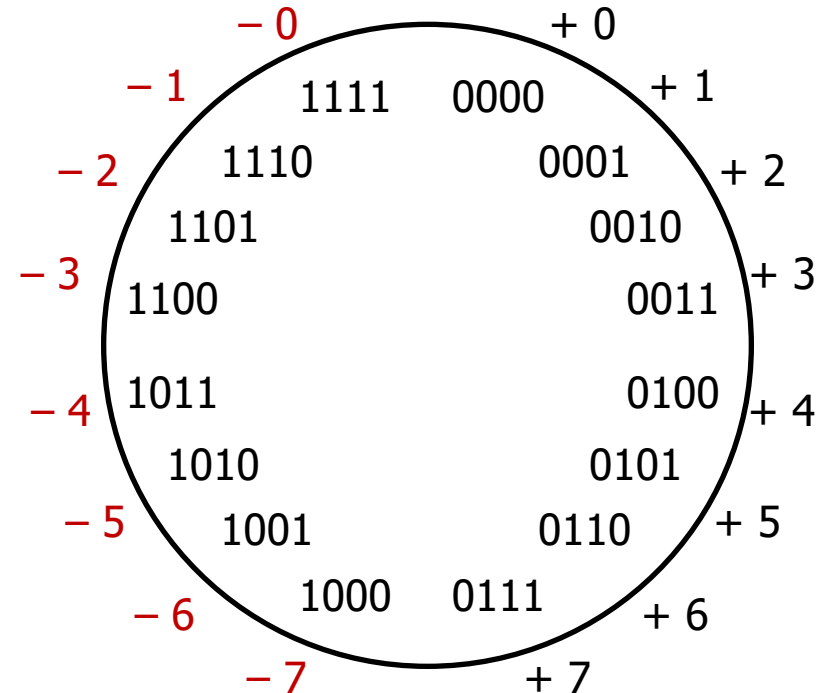


- Negatives “increment” in wrong direction!



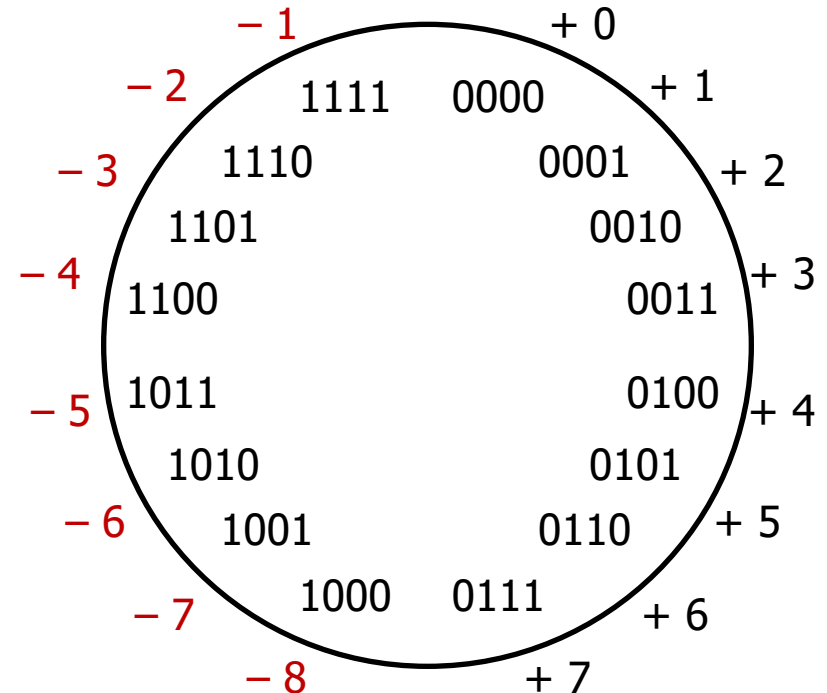
# 3) Two's Complement

- Let's fix these problems:
  - “Flip” negative encodings so incrementing works



# 3) Two's Complement

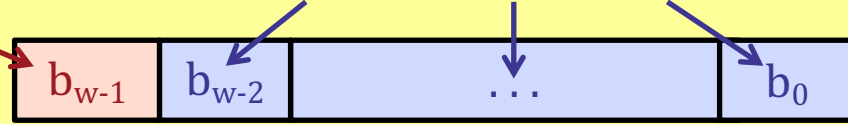
- Let's fix these problems:
  - 1) “Flip” negative encodings so incrementing works
  - 2) “Shift” negative numbers to eliminate  $-0$
- MSB *still* indicates sign!
  - This is why we represent one more negative than positive number ( $-2^{N-1}$  to  $2^{N-1} - 1$ )



# Two's Complement Negatives

- Accomplished with one neat mathematical trick!

$b_{w-1}$  has weight  $-2^{w-1}$ , other bits have usual weights  $+2^i$



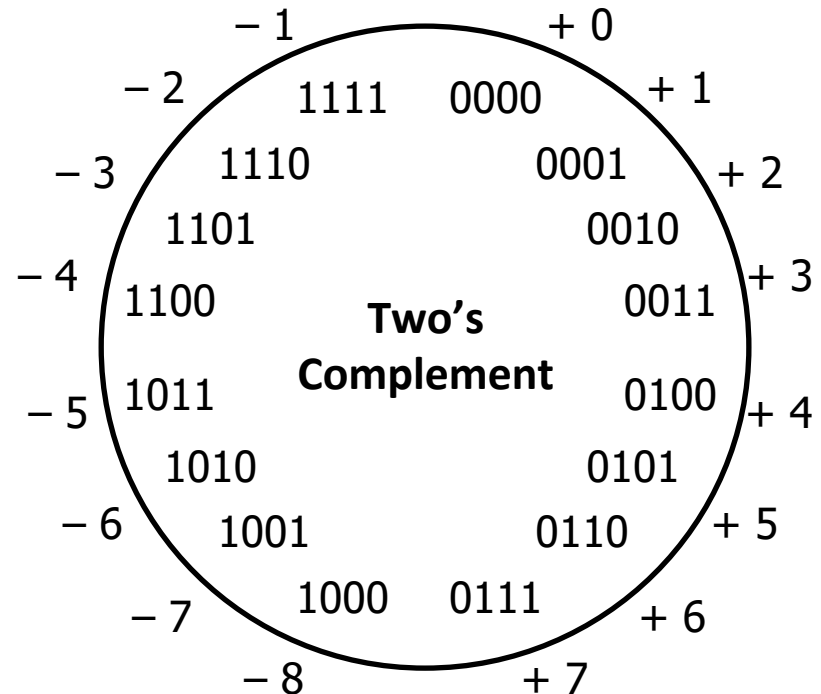
- 4-bit Examples:

- $1010_2$  unsigned:  
 $1*2^3 + 0*2^2 + 1*2^1 + 0*2^0 = 10$
- $1010_2$  two's complement:  
 $-1*2^3 + 0*2^2 + 1*2^1 + 0*2^0 = -6$

- 1 represented as:

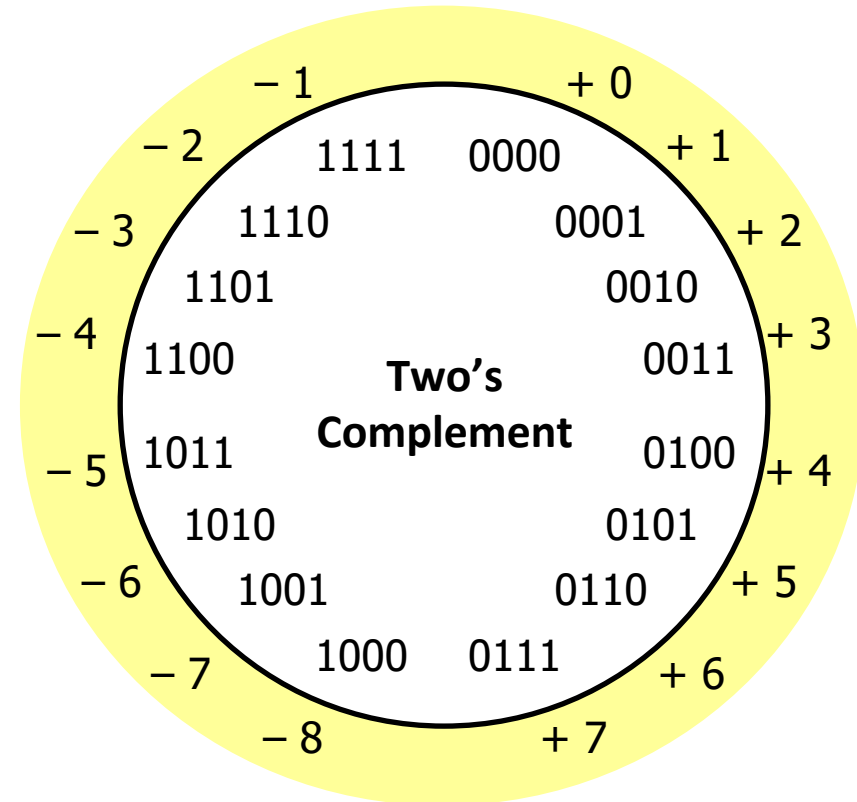
$$1111_2 = -2^3 + (2^3 - 1)$$

- MSB makes it super negative, add up all the other bits to get back up to -1



# Why Two's Complement is So Great

- Roughly same number of (+) and (−) numbers
- Positive number encodings match unsigned
- Single zero
- All zeros encoding = 0
- Simple negation procedure:
  - Get negative representation of any integer by taking bitwise complement and then adding one!  
(  $\sim x + 1 == -x$  )



# Summary

- Bit-level operators allow for fine-grained manipulations of data
  - Bitwise AND ( $\&$ ), OR ( $|$ ), and NOT ( $\sim$ ) different than logical AND ( $\&\&$ ), OR ( $||$ ), and NOT ( $!$ )
  - Especially useful with bit masks
- Choice of *encoding scheme* is important
  - Tradeoffs based on size requirements and desired operations
- Integers represented using unsigned and two's complement representations
  - Limited by fixed bit width
  - We'll examine arithmetic operations in the next lecture