

Lecture 02

Bits, Bytes and Integers – part 2

Euhyun Moon, Ph.D.
Machine Learning Systems (MLSys) Lab
Computer Science and Engineering
Sogang University



Slides adapted from Randy Bryant and Dave O'Hallaron: Introduction to Computer Systems, CMU

Two's Complement Arithmetic

- The same addition procedure works for both unsigned and two's complement integers
 - **Simplifies hardware:** only one algorithm for addition
 - **Algorithm:** simple addition, **discard the highest carry bit**
 - Called modular addition: result is sum *modulo* 2^w
- **4-bit Examples:**

HW	TC
0100 +0011	
=	

HW	TC
1100 +0011	
=	

HW	TC
0100 +1101	
=	

Why Does Two's Complement Work?

- For all representable positive integers x , we want:

$$\begin{array}{r} \textit{bit representation of } x \\ + \textit{ bit representation of } -x \\ \hline 0 \end{array} \quad (\text{ignoring the carry-out bit})$$

- What are the 8-bit negative encodings for the following?

$$\begin{array}{r} 00000001 \\ + \text{????????} \\ \hline 00000000 \end{array}$$

$$\begin{array}{r} 00000010 \\ + \text{????????} \\ \hline 00000000 \end{array}$$

$$\begin{array}{r} 11000011 \\ + \text{????????} \\ \hline 00000000 \end{array}$$

Why Does Two's Complement Work?

- For all representable positive integers x , we want:

$$\begin{array}{r} \textit{bit representation of } x \\ + \textit{ bit representation of } -x \\ \hline 0 \end{array} \quad (\text{ignoring the carry-out bit})$$

- What are the 8-bit negative encodings for the following?

$$\begin{array}{r} 00000001 \\ + 11111111 \\ \hline 100000000 \end{array}$$

$$\begin{array}{r} 00000010 \\ + 11111110 \\ \hline 100000000 \end{array}$$

$$\begin{array}{r} 11000011 \\ + 00111101 \\ \hline 100000000 \end{array}$$

These are the bitwise complement plus 1!

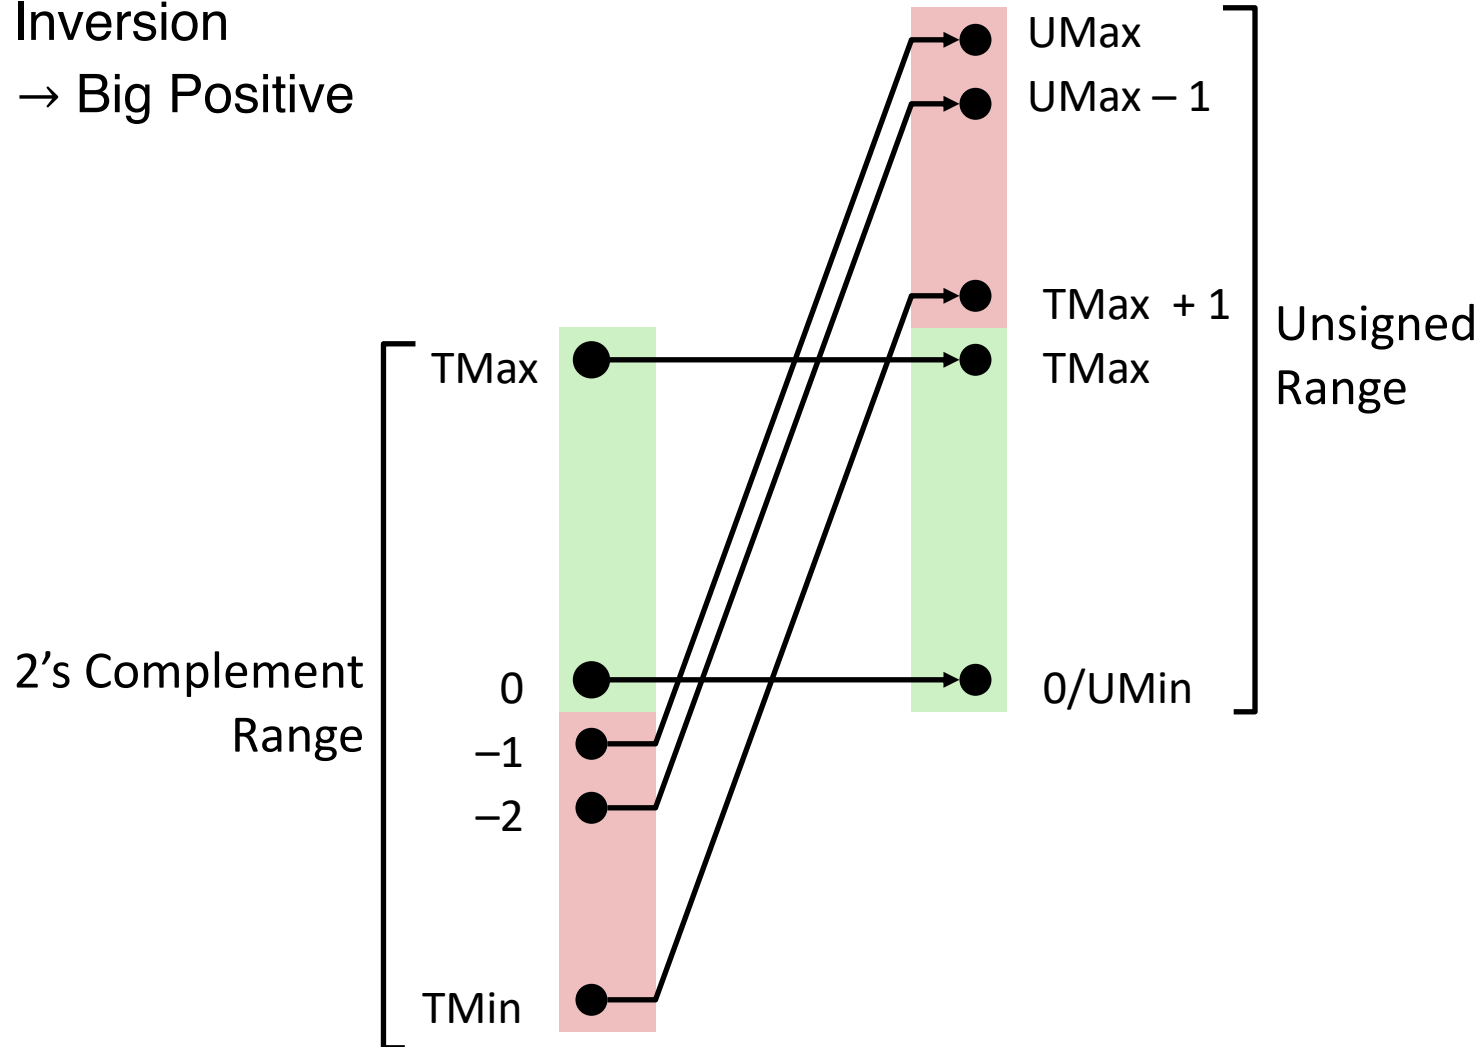
$$-x == \sim x + 1$$

Integers

- **Binary representation of integers**
 - **Unsigned and signed**
 - **Casting in C**
- Consequences of finite width representations
 - Sign extension, overflow
- Shifting and arithmetic operations

Signed/Unsigned Conversion Visualized

- Two's Complement \rightarrow Unsigned
 - Ordering Inversion
 - Negative \rightarrow Big Positive



Values To Remember

- Unsigned Values

- UMin = 0b00...0
= 0
- UMax = 0b11...1
= $2^w - 1$

- Two's Complement Values

- TMin = 0b10...0
= -2^{w-1}
- TMax = 0b01...1
= $2^{w-1} - 1$
- -1 = 0b11...1

- **Example:** Values for $w = 64$

	Decimal	Hex
UMax	18,446,744,073,709,551,615	FF FF FF FF FF FF FF FF
TMax	9,223,372,036,854,775,807	7F FF FF FF FF FF FF FF
TMin	-9,223,372,036,854,775,808	80 00 00 00 00 00 00 00
-1	-1	FF FF FF FF FF FF FF FF
0	0	00 00 00 00 00 00 00 00

In C: Signed vs. Unsigned

- Casting
 - Bits are unchanged, just interpreted differently!
 - `int tx, ty;`
 - `unsigned int ux, uy;`
 - *Explicit* casting
 - `tx = (int) ux;`
 - `uy = (unsigned int) ty;`
 - *Implicit* casting can occur during assignments or function calls
 - `tx = ux;`
 - `uy = ty;`

Casting Surprises!

- Integer literals (constants)
 - By default, integer constants are considered *signed* integers
 - Hex constants already have an explicit binary representation
 - Use “U” (or “u”) suffix to explicitly force *unsigned*
 - **Examples:** 0U, 4294967259u
- Expression Evaluation
 - When you mixed unsigned and signed in a single expression, then **signed values are implicitly cast to unsigned**
 - Including comparison operators <, >, ==, <=, >=

Integers

- Binary representation of integers
 - Unsigned and signed
 - Casting in C
- **Consequences of finite width representations**
 - **Sign extension, overflow**
- Shifting and arithmetic operations

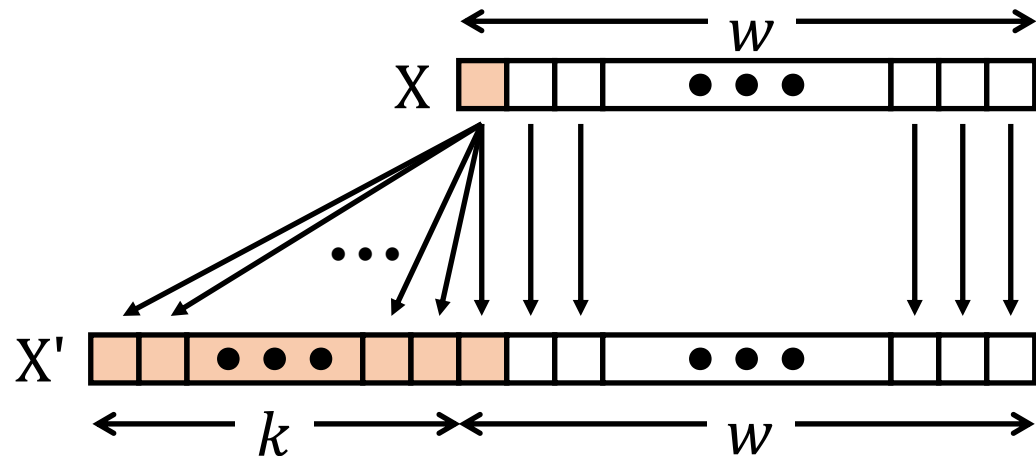
Sign Extension

- What happens if you convert a *signed* integral data type to a larger one?
 - *e.g.* char → short → int → long
 - **4-bit → 8-bit Example:**
 - Positive Case
 - ✓ • Add 0's?
- 4-bit:** 0010 = +2

8-bit: 00000010 = +2
- Negative Case?

Sign Extension

- **Task:** Given a w -bit signed integer X , convert it to $w+k$ -bit signed integer X' *with the same value*
- **Rule:** Add k copies of sign bit
 - Let x_i be the i -th digit of X in binary
 - $X' = \underbrace{x_{w-1}, \dots, x_{w-1}}_{k \text{ copies of MSB}}, \underbrace{x_{w-1}, x_{w-2}, \dots, x_1, x_0}_{\text{original } X}$



Sign Extension Example

- Convert from smaller to larger integral data types
- C automatically performs sign extension
 - Java too

```
short int x = 12345;  
int     ix = (int) x;  
short int y = -12345;  
int     iy = (int) y;
```

Var	Decimal	Hex	Binary
x	12345	30 39	00110000 00111001
ix	12345	00 00 30 39	00000000 00000000 00110000 00111001
y	-12345	CF C7	11001111 11000111
iy	-12345	FF FF CF C7	11111111 11111111 11001111 11000111

Arithmetic Overflow

Bits	Unsigned	Signed
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	-8
1001	9	-7
1010	10	-6
1011	11	-5
1100	12	-4
1101	13	-3
1110	14	-2
1111	15	-1

- When a calculation produces a result that can't be represented in the current encoding scheme
 - Integer range limited by fixed width
 - Can occur in both the positive and negative directions
- C and Java ignore overflow exceptions
 - You end up with a bad value in your program and no warning/indication... oops!

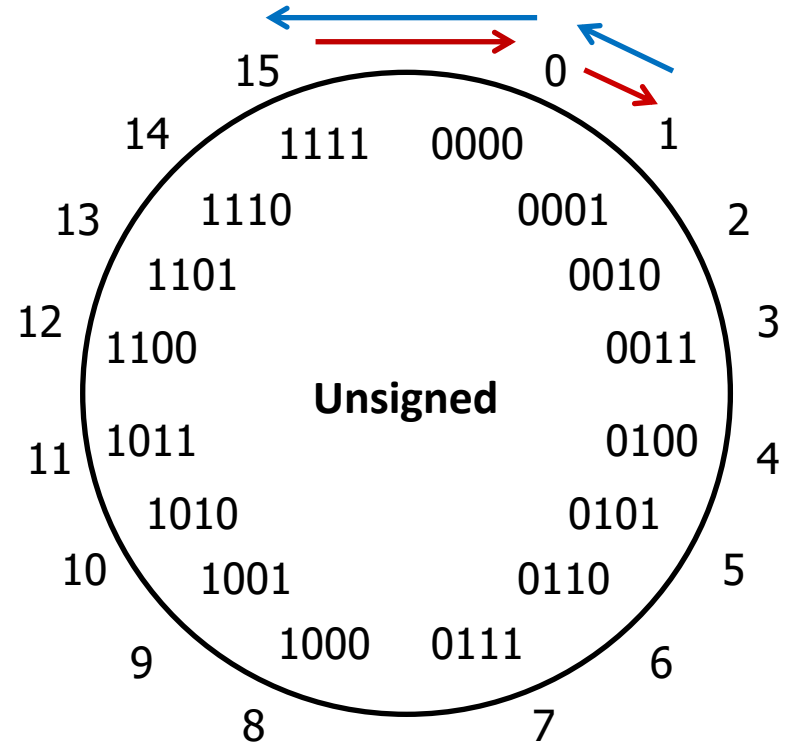
Overflow: Unsigned

- **Addition:** drop carry bit (-2^N)

15	1111
+ 2	+ 0010
17	1 0001
1	

- **Subtraction:** borrow ($+2^N$)

1	10001
- 2	- 0010
-1	1111
15	



$\pm 2^N$ because of
modular arithmetic

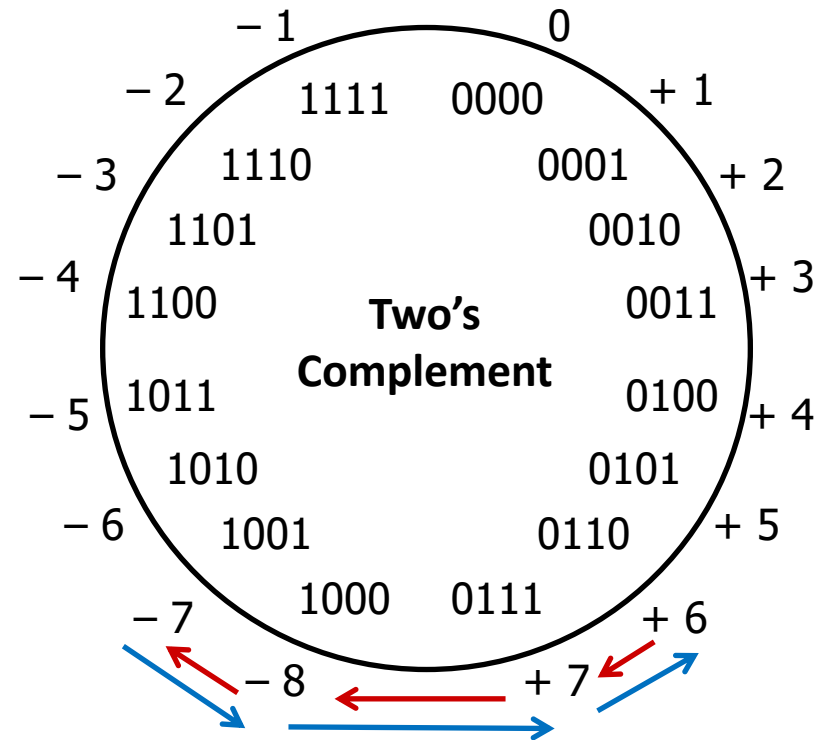
Overflow: Two's Complement

- **Addition:** $(+) + (+) = (-)$ result?

$$\begin{array}{r} 6 \\ + 3 \\ \hline \cancel{9} \\ -7 \end{array} \qquad \begin{array}{r} 0110 \\ + 0011 \\ \hline 1001 \end{array}$$

- **Subtraction:** $(-) + (-) = (+)$?

$$\begin{array}{r} -7 \\ - 3 \\ \hline \cancel{-10} \\ 6 \end{array} \qquad \begin{array}{r} 1001 \\ - 0011 \\ \hline 0110 \end{array}$$



For signed: overflow if operands have same sign and result's sign is different

Integers

- Binary representation of integers
 - Unsigned and signed
 - Casting in C
- Consequences of finite width representations
 - Sign extension, overflow
- **Shifting and arithmetic operations**

Shift Operations

- Throw away (drop) extra bits that “fall off” the end
- Left shift ($x \ll n$) bit vector x by n positions
 - Fill with 0's on right
- Right shift ($x \gg n$) bit-vector x by n positions
 - Logical shift (for **unsigned** values)
 - Fill with 0's on left
 - Arithmetic shift (for **signed** values)
 - Replicate most significant bit on left (maintains sign of x)

Shift Operations

- Left shift ($x \ll n$)
 - Fill with 0's on right

- Right shift ($x \gg n$)
 - Logical shift (for **unsigned** values)
 - Fill with 0's on left
 - Arithmetic shift (for **signed** values)
 - Replicate most significant bit on left

- Notes:
 - Shifts by $n < 0$ or $n \geq w$ (w is bit width of x) are *undefined*
 - **In C:** behavior of \gg is determined by compiler
 - In gcc / C lang, depends on data type of x (signed/unsigned)
 - **In Java:** logical shift is \ggg and arithmetic shift is \gg

	x	0010	0010
	$x \ll 3$	0001	0 000
logical:	$x \ggg 2$	00 00	1000
arithmetic:	$x \gg 2$	00 00	1000

	x	1010	0010
	$x \ll 3$	0001	0 000
logical:	$x \ggg 2$	00 10	1000
arithmetic:	$x \gg 2$	11 10	1000

Shifting Arithmetic?

- What are the following computing?
 - $x \gg n$
 - $0b\ 0100 \gg 1 = 0b\ 0010$
 - $0b\ 0100 \gg 2 = 0b\ 0001$
 - Divide by 2^n
 - $x \ll n$
 - $0b\ 0001 \ll 1 = 0b\ 0010$
 - $0b\ 0001 \ll 2 = 0b\ 0100$
 - Multiply by 2^n

Shifting is faster than general multiply and divide operations!

Left Shifting Arithmetic 8-bit Example

- No difference in left shift operation for unsigned and signed numbers (just manipulates bits)
 - Difference comes during interpretation: $x * 2^n$?

		Signed	Unsigned
$x = 25;$	00011001 =	25	25
$L1 = x \ll 2;$	0001100100 =	100	100
$L2 = x \ll 3;$	00011001000 =	-56	200
$L3 = x \ll 4;$	000110010000 =	-112	144

signed overflow

unsigned overflow

Right Shifting Arithmetic 8-bit Examples

- **Reminder:** C operator `>>` does *logical* shift on **unsigned** values and *arithmetic* shift on **signed** values
 - Logical Shift: $x / 2^n$?

`xu = 240u;` 11110000 = 240

`R1u=xu>>3;` 00011110000 = 30

`R2u=xu>>5;` 0000011110000 = 7

rounding (down)

Right Shifting Arithmetic 8-bit Examples

- **Reminder:** C operator `>>` does *logical* shift on **unsigned** values and *arithmetic* shift on **signed** values
 - **Arithmetic** Shift: $x / 2^n$?

`xs = -16;` 11110000 = -16

`R1s=xs>>3;` 11111110000 = -2

`R2s=xs>>5;` 1111111110000 = -1

rounding (down)

Summary

- Sign and unsigned variables in C
 - Bit pattern remains the same, just *interpreted* differently
 - Strange things can happen with our arithmetic when we convert/cast between sign and unsigned numbers
 - Type of variables affects behavior of operators (shifting, comparison)
- We can only represent so many numbers in w bits
 - When we exceed the limits, *arithmetic overflow* occurs
 - *Sign extension* tries to preserve value when expanding
- Shifting is a useful bitwise operator
 - Right shifting can be arithmetic (sign) or logical (0)
 - Can be used in multiplication with constant or bit masking