

Lecture 03

Memory, Data and Addressing – part 2

Euhyun Moon, Ph.D.

Machine Learning Systems (MLSys) Lab

Computer Science and Engineering

Sogang University



Slides adapted from Randy Bryant and Dave O'Hallaron: Introduction to Computer Systems, CMU

Memory, Data, and Addressing

- Hardware - High Level Overview
- Representing information as bits and bytes
 - Memory is a byte-addressable array
 - Machine “word” size = address size = register size
 - Endianness – ordering bytes in memory
- **Manipulating data in memory using C**
 - **Assignment**
 - **Pointers, pointer arithmetic, and arrays**
- Boolean algebra and bit-level manipulations

Addresses and Pointers in C

- `&` = “address of” operator
- `*` = “value at address” or “dereference” operator

`*` is also used with variable declarations

```
int* ptr;
```

Declares a variable, `ptr`, that is a pointer to (i.e. holds the address of) an `int` in memory

```
int x = 5;
```

```
int y = 2;
```

Declares two variables, `x` and `y`, that hold `ints`, and *initializes* them to 5 and 2, respectively

```
ptr = &x;
```

Sets `ptr` to the address of `x` (“`ptr` points to `x`”)

```
y = 1 + *ptr;
```

“Dereference `ptr`”

Sets `y` to “1 plus the value stored at the address held by `ptr`.” Because `ptr` points to `x`, this is equivalent to `y=1+x`;

What is `*(&y)` ?

Assignment in C

- A variable is represented by a location
- Declaration \neq initialization (initially holds random data)
- `int x, y;`
 - `x` is at address `0x04`, `y` is at `0x18`

	0x00	0x01	0x02	0x03	
0x00	A7	00	32	00	
0x04	00	01	29	F3	x
0x08	EE	EE	EE	EE	
0x0C	FA	CE	CA	FE	
0x10	26	00	00	00	
0x14	00	00	10	00	
0x18	01	00	00	00	y
0x1C	FF	00	F4	96	
0x20	DE	AD	BE	EF	
0x24	00	00	00	00	

Assignment in C

32-bit example
(pointers are 32-bits wide)

little-endian

- A variable is represented by a location
- Declaration \neq initialization (initially holds random data)
- `int x, y;`
 - `x` is at address `0x04`, `y` is at `0x18`

	0x00	0x01	0x02	0x03	
0x00					
0x04	00	01	29	F3	x
0x08					
0x0C					
0x10					
0x14					
0x18	01	00	00	00	y
0x1C					
0x20					
0x24					

Assignment in C

32-bit example
(pointers are 32-bits wide)

- left-hand side = right-hand side;
 - LHS must evaluate to a *location*
 - RHS must evaluate to a *value* (could be an address)
 - Store RHS value at LHS location
- `int x, y;`
- `x = 0;`

& = “address of”
* = “dereference”

	0x00	0x01	0x02	0x03	
0x00					
0x04	00	00	00	00	x
0x08					
0x0C					
0x10					
0x14					
0x18	01	00	00	00	y
0x1C					
0x20					
0x24					

Assignment in C

32-bit example
(pointers are 32-bits wide)

- left-hand side = right-hand side;
 - LHS must evaluate to a *location*
 - RHS must evaluate to a *value* (could be an address)
 - Store RHS value at LHS location
- `int x, y;`
- `x = 0;`
- `y = 0x3CD02700;`

	0x00	0x01	0x02	0x03	
0x00					
0x04	00	00	00	00	x
0x08					
0x0C					
0x10					
0x14					
0x18	00	27	D0	3C	y
0x1C					
0x20					
0x24					

little endian!

Assignment in C

32-bit example
(pointers are 32-bits wide)

- left-hand side = right-hand side;
 - LHS must evaluate to a *location*
 - RHS must evaluate to a *value* (could be an address)
 - Store RHS value at LHS location
- `int x, y;`
- `x = 0;`
- `y = 0x3CD02700;`
- `x = y + 3;`
 - Get value at `y`, add 3, store in `x`

& = “address of”
* = “dereference”

	0x00	0x01	0x02	0x03	
0x00					x
0x04	03	27	D0	3C	
0x08					
0x0C					
0x10					y
0x14					
0x18	00	27	D0	3C	
0x1C					
0x20					
0x24					

Assignment in C

32-bit example
(pointers are 32-bits wide)

- left-hand side = right-hand side;
 - LHS must evaluate to a *location*
 - RHS must evaluate to a *value* (could be an address)
 - Store RHS value at LHS location
- **int** x, y;
- x = 0;
- y = 0x3CD02700;
- x = y + 3;
 - Get value at y, add 3, store in x
- **int*** z;
 - z is at address 0x20

& = “address of”
* = “dereference”

	0x00	0x01	0x02	0x03	
0x00					
0x04	03	27	D0	3C	x
0x08					
0x0C					
0x10					
0x14					
0x18	00	27	D0	3C	y
0x1C					
0x20	DE	AD	BE	EF	z
0x24					

Assignment in C

32-bit example
(pointers are 32-bits wide)

- left-hand side = right-hand side;
 - LHS must evaluate to a *location*
 - RHS must evaluate to a *value* (could be an address)
 - Store RHS value at LHS location
- **int** x, y;
- x = 0;
- y = 0x3CD02700;
- x = y + 3;
 - Get value at y, add 3, store in x
- **int*** z = &y + 3;
 - Get address of y, “add 3”, store in z

& = “address of”
* = “dereference”

	0x00	0x01	0x02	0x03	
0x00					
0x04	03	27	D0	3C	x
0x08					
0x0C					
0x10					
0x14					
0x18	00	27	D0	3C	y
0x1C					
0x20	24	00	00	00	z
0x24					

Pointer arithmetic

Pointer Arithmetic

- Pointer arithmetic is scaled by the size of target type
 - In this example, `sizeof(int) = 4`
- `int* z = &y + 3;`
 - Get address of `y`, add `3*sizeof(int)`, store in `z`
 - $\&y = 0x18 = 1*16^1 + 8*16^0 = 24$
 - $24 + 3*(4) = 36 = 2*16^1 + 4*16^0 = 0x24$
- **Pointer arithmetic can be dangerous!**
 - Can easily lead to bad memory accesses
 - Be careful with data types and *casting*

Assignment in C

- `int x, y;`
- `x = 0;`
- `y = 0x3CD02700;`
- `x = y + 3;`
 - Get value at `y`, add 3, store in `x`
- `int* z = &y + 3;`
 - Get address of `y`, add **12**, store in `z`
- `*z = y;`
 - What does this do?

32-bit example
(pointers are 32-bits wide)

& = “address of”
* = “dereference”

	0x00	0x01	0x02	0x03	
0x00					
0x04	03	27	D0	3C	x
0x08					
0x0C					
0x10					
0x14					
0x18	00	27	D0	3C	y
0x1C					
0x20	24	00	00	00	z
0x24					

Assignment in C

- `int x, y;`
- `x = 0;`
- `y = 0x3CD02700;`
- `x = y + 3;`
 - Get value at `y`, add 3, store in `x`
- `int* z = &y + 3;`
 - Get address of `y`, add **12**, store in `z`

The target of a pointer
is also a location

- `*z = y;`
 - Get value of `y`, put in address stored in `z`

32-bit example
(pointers are 32-bits wide)

`&` = "address of"
`*` = "dereference"

	0x00	0x01	0x02	0x03	
0x00					
0x04	03	27	D0	3C	x
0x08					
0x0C					
0x10					
0x14					
0x18	00	27	D0	3C	y
0x1C					
0x20	24	00	00	00	z
0x24	00	27	D0	3C	

Arrays in C

Arrays are adjacent locations in memory storing the same type of data object

`a` (array name) returns the array's address

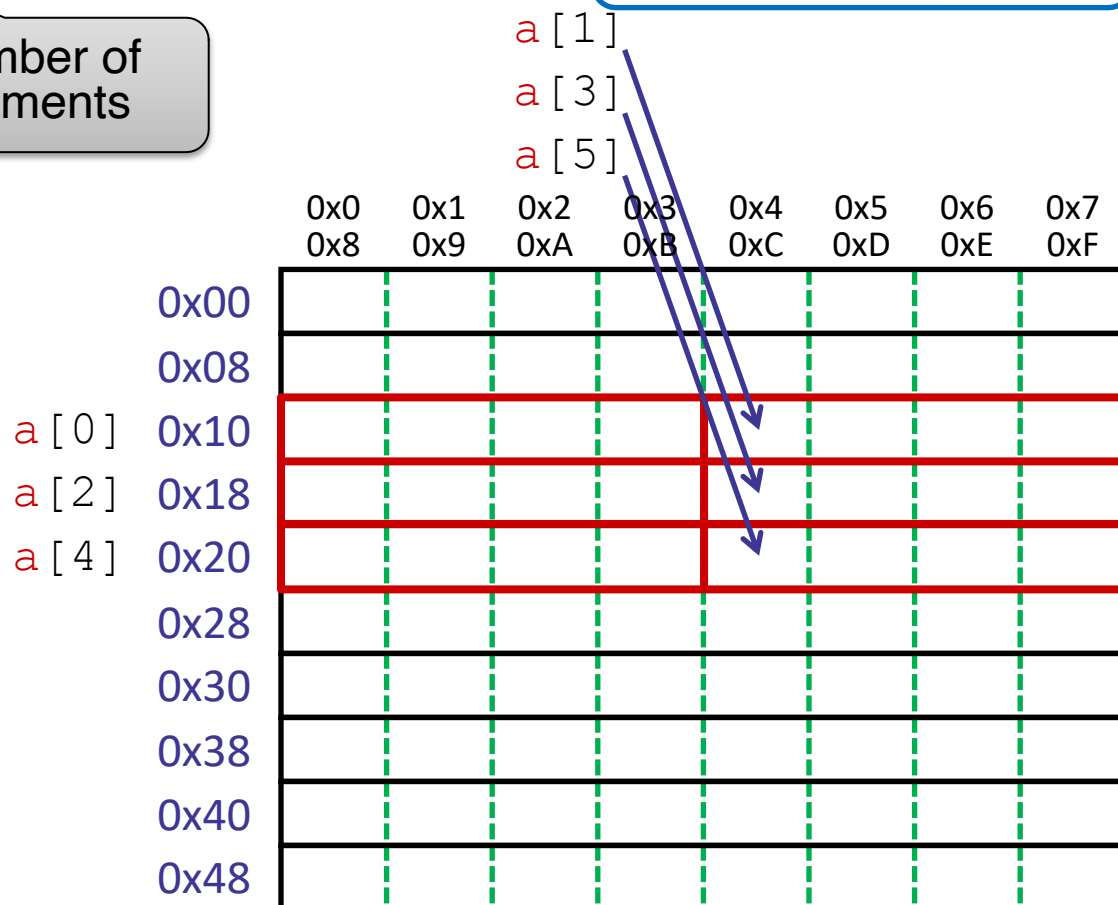
Declaration: `int a[6];`

element type

name

number of elements

64-bit example
(pointers are 64-bits wide)



Arrays in C

Arrays are adjacent locations in memory storing the same type of data object

`a` (array name) returns the array's address

`&a[i]` is the address of `a[0]` plus `i` times the element size in bytes

Declaration: `int a[6];`

Indexing: `a[0] = 0x015f;`
`a[5] = a[0];`

		0x0	0x1	0x2	0x3	0x4	0x5	0x6	0x7
		0x8	0x9	0xA	0xB	0xC	0xD	0xE	0xF
0x00									
0x08									
<code>a[0]</code>	0x10	5F	01	00	00				
<code>a[2]</code>	0x18								
<code>a[4]</code>	0x20					5F	01	00	00
0x28									
0x30									
0x38									
0x40									
0x48									

Arrays in C

Arrays are adjacent locations in memory storing the same type of data object

`a` (array name) returns the array's address

`&a[i]` is the address of `a[0]` plus `i` times the element size in bytes

Declaration: `int a[6];`

Indexing: `a[0] = 0x015f;`
`a[5] = a[0];`

No bounds checking: `a[6] = 0xBAD;`
`a[-1] = 0xBAD;`

	0x0 0x8	0x1 0x9	0x2 0xA	0x3 0xB	0x4 0xC	0x5 0xD	0x6 0xE	0x7 0xF
0x00								
0x08					AD	0B	00	00
<code>a[0]</code> 0x10	5F	01	00	00				
<code>a[2]</code> 0x18								
<code>a[4]</code> 0x20					5F	01	00	00
0x28	AD	0B	00	00				
0x30								
0x38								
0x40								
0x48								

Arrays in C

Arrays are adjacent locations in memory storing the same type of data object

`a` (array name) returns the array's address

`&a[i]` is the address of `a[0]` plus `i` times the element size in bytes

Declaration: `int a[6];`

Indexing: `a[0] = 0x015f;`
`a[5] = a[0];`

No bounds checking: `a[6] = 0xBAD;`
`a[-1] = 0xBAD;`

Pointers: `int* p;`

equivalent $\left\{ \begin{array}{l} p = a; \\ p = \&a[0]; \\ *p = 0xA; \end{array} \right.$

`a[0]`
`a[2]`
`a[4]`

`p`

	0x0 0x8	0x1 0x9	0x2 0xA	0x3 0xB	0x4 0xC	0x5 0xD	0x6 0xE	0x7 0xF
0x00								
0x08					AD	0B	00	00
0x10	0A	00	00	00				
0x18								
0x20					5F	01	00	00
0x28	AD	0B	00	00				
0x30								
0x38								
0x40	10	00	00	00	00	00	00	00
0x48								

Arrays in C

Arrays are adjacent locations in memory storing the same type of data object

`a` (array name) returns the array's address

`&a[i]` is the address of `a[0]` plus `i` times the element size in bytes

Declaration: `int a[6];`

Indexing: `a[0] = 0x015f;`
`a[5] = a[0];`

No bounds checking: `a[6] = 0xBAD;`
`a[-1] = 0xBAD;`

Pointers: `int* p;`

equivalent $\begin{cases} p = a; \\ p = \&a[0]; \\ *p = 0xA; \end{cases}$

`a[0]`
`a[2]`
`a[4]`

array indexing = address arithmetic
(both scaled by the size of the type)

equivalent $\begin{cases} p[1] = 0xB; \\ *(p+1) = 0xB; \\ p = p + 2; \end{cases}$

`p`

	0x0 0x8	0x1 0x9	0x2 0xA	0x3 0xB	0x4 0xC	0x5 0xD	0x6 0xE	0x7 0xF
0x00								
0x08					AD	0B	00	00
0x10	0A	00	00	00	0B	00	00	00
0x18								
0x20					5F	01	00	00
0x28	AD	0B	00	00				
0x30								
0x38								
0x40	10	00	00	00	00	00	00	00
0x48								

Arrays in C

Arrays are adjacent locations in memory storing the same type of data object

`a` (array name) returns the array's address

`&a[i]` is the address of `a[0]` plus `i` times the element size in bytes

Declaration: `int a[6];`

Indexing: `a[0] = 0x015f;`
`a[5] = a[0];`

No bounds checking: `a[6] = 0xBAD;`
`a[-1] = 0xBAD;`

Pointers: `int* p;`

equivalent $\left\{ \begin{array}{l} p = a; \\ p = \&a[0]; \\ *p = 0xA; \end{array} \right.$

`a[0]`
`a[2]`
`a[4]`

array indexing = address arithmetic
 (both scaled by the size of the type)

equivalent $\left\{ \begin{array}{l} p[1] = 0xB; \\ *(p+1) = 0xB; \\ p = p + 2; \\ *p = a[1] + 1; \end{array} \right.$

`p`

	0x0	0x1	0x2	0x3	0x4	0x5	0x6	0x7
	0x8	0x9	0xA	0xB	0xC	0xD	0xE	0xF
0x00								
0x08					AD	0B	00	00
0x10	0A	00	00	00	0B	00	00	00
0x18	0C	00	00	00				
0x20					5F	01	00	00
0x28	AD	0B	00	00				
0x30								
0x38								
0x40	18	00	00	00	00	00	00	00
0x48								

Representing strings

- C-style string stored as an array of bytes (**char***)
 - No “String” keyword, unlike Java
 - Elements are one-byte **ASCII codes** for each character

32	space	48	0	64	@	80	P	96	`	112	p
33	!	49	1	65	A	81	Q	97	a	113	q
34	”	50	2	66	B	82	R	98	b	114	r
35	#	51	3	67	C	83	S	99	c	115	s
36	\$	52	4	68	D	84	T	100	d	116	t
37	%	53	5	69	E	85	U	101	e	117	u
38	&	54	6	70	F	86	V	102	f	118	v
39	'	55	7	71	G	87	W	103	g	119	w
40	(56	8	72	H	88	X	104	h	120	x
41)	57	9	73	I	89	Y	105	i	121	y
42	*	58	:	74	J	90	Z	106	j	122	z
43	+	59	;	75	K	91	[107	k	123	{
44	,	60	<	76	L	92	\	108	l	124	
45	-	61	=	77	M	93]	109	m	125	}
46	.	62	>	78	N	94	^	110	n	126	~
47	/	63	?	79	O	95	_	111	o	127	del

ASCII: American Standard Code for Information Interchange

Null-Terminated Strings

- **Example:** "Heung-min Son" stored as a 14-byte array

Decimal:	72	101	117	110	103	45	109	105	110	32	83	111	110	0
Hex:	0x48	0x65	0x75	0x6E	0x67	0x2D	0x6D	0x69	0x6E	0x20	0x53	0x6F	0x6E	0x00
Text:	H	e	u	n	g	-	m	i	n		S	o	n	\0

- Last character followed by a 0 byte (' \0 ')
(a.k.a. "null terminator")
 - Must take into account when allocating space in memory
 - Note that ' 0 ' \neq ' \0 ' (*i.e.* character 0 has non-zero value)
- How do we compute the length of a string?
 - Traverse array until null terminator encountered

Endianness and Strings

C (char = 1 byte)

```
char s[6] = "12345";
```

String literal

0x31 = 49 decimal = ASCII '1'

IA32, x86-64

(little-endian)

SPARC

(big-endian)

0x00	31	↔	31	0x00	'1'
0x01	32	↔	32	0x01	'2'
0x02	33	↔	33	0x02	'3'
0x03	34	↔	34	0x03	'4'
0x04	35	↔	35	0x04	'5'
0x05	00	↔	00	0x05	'\0'

- Byte ordering (endianness) is not an issue for 1-byte values
 - The whole array does not constitute a single value
 - Individual elements are values; chars are single bytes

Examining Data Representations

- Code to print byte representation of data
 - Any data type can be treated as a *byte array* by **casting** it to `char`
 - C has **unchecked** casts **!! DANGER !!**

```
void show_bytes(char* start, int len) {  
    int i;  
    for (i = 0; i < len; i++)  
        printf("%p\t0x%.2hhX\n", start+i, *(start+i));  
    printf("\n");  
}
```

`printf` **directives:**

<code>%p</code>	Print pointer
<code>\t</code>	Tab
<code>%.2hhx</code>	Print value as char (hh) in hex (X), padding to 2 digits (.2)
<code>\n</code>	New line

Examining Data Representations

- Code to print byte representation of data
 - Any data type can be treated as a *byte array* by **casting** it to `char`
 - C has **unchecked** casts **!! DANGER !!**

```
void show_bytes(char* start, int len) {
    int i;
    for (i = 0; i < len; i++)
        printf("%p\t0x%.2hhX\n", start+i, *(start+i));
    printf("\n");
}
```

```
void show_int(int x) {
    show_bytes((char *) &x, sizeof(int));
}
```


show_bytes Execution Example

```
int x = 12345; // 0x00 00 30 39
printf("int x = %d;\n", x);
show_int(x);    // show_bytes((char *) &x, sizeof(int));
```

- Result (Linux x86-64):

- **Note:** The addresses will change on each run (try it!), but fall in same general range

```
int x = 12345;
0x7fffb7f71dbc      0x39
0x7fffb7f71dbd      0x30
0x7fffb7f71dbe      0x00
0x7fffb7f71dbf      0x00
```

Summary

- Assignment in C results in value being put in memory location
- Pointer is a C representation of a data address
 - `&` = “address of” operator
 - `*` = “value at address” or “dereference” operator
- Pointer arithmetic scales by size of target type
 - Convenient when accessing array-like structures in memory
 - Be careful when using – particularly when *casting* variables
- Arrays are adjacent locations in memory storing the same type of data object
 - Strings are null-terminated arrays of characters (ASCII)