

# Lecture 04

## Machine-Level Programming – part 1

Euhyun Moon, Ph.D.  
Machine Learning Systems (MLSys) Lab  
Computer Science and Engineering  
Sogang University




Slides adapted from Randy Bryant and Dave O'Hallaron: Introduction to Computer Systems, CMU

# Definitions

- **Architecture (ISA):** The parts of a processor design that one needs to understand to write assembly code
  - “What is directly visible to software”
  - The “contract” or “blueprint” between hardware and software
- **Microarchitecture:** Implementation of the architecture

# Instruction Set Architectures

- The ISA defines:
  - The system's **state** (e.g., registers, memory, program counter)
  - The **instructions** the CPU can execute
  - The **effect** that each of these instructions will have on the system state




# General ISA Design Decisions

- Instructions
  - What instructions are available? What do they do?
  - How are they encoded?
- Registers
  - How many registers are there?
  - How wide are they?
- Memory
  - How do you specify a memory location?

# Instruction Set Philosophies

- *Complex Instruction Set Computing (CISC):*  
Add more and more elaborate and specialized instructions as needed
  - Lots of tools for programmers to use, but hardware must be able to handle all instructions
  - x86-64 is CISC, but only a small subset of instructions encountered with Linux programs
- *Reduced Instruction Set Computing (RISC):*  
Keep instruction set small and regular
  - Easier to build fast hardware
  - Let software do the complicated operations by composing simpler ones

# Mainstream ISAs



## x86

<b>Designer</b>	Intel, AMD
<b>Bits</b>	16-bit, 32-bit and 64-bit
<b>Introduced</b>	1978 (16-bit), 1985 (32-bit), 2003 (64-bit)
<b>Design</b>	CISC
<b>Type</b>	Register-memory
<b>Encoding</b>	Variable (1 to 15 bytes)
<b>Endianness</b>	Little

Macbooks & PCs  
(Core i3, i5, i7, M)  
[x86-64 Instruction Set](#)



## ARM architectures

<b>Designer</b>	ARM Holdings
<b>Bits</b>	32-bit, 64-bit
<b>Introduced</b>	1985; 31 years ago
<b>Design</b>	RISC
<b>Type</b>	Register-Register
<b>Encoding</b>	AArch64/A64 and AArch32/A32 use 32-bit instructions, T32 (Thumb-2) uses mixed 16- and 32-bit instructions. ARMv7 <a href="#">user-space compatibility</a> <sup>[1]</sup>
<b>Endianness</b>	Bi (little as default)

Smartphone-like devices  
(iPhone, iPad, Raspberry Pi)  
[ARM Instruction Set](#)



## RISC-V

<b>Designer</b>	University of California, Berkeley
<b>Bits</b>	32 · 64 · 128
<b>Introduced</b>	2010
<b>Design</b>	RISC
<b>Type</b>	Load-store
<b>Encoding</b>	Variable
<b>Endianness</b>	Little <sup>[1][3]</sup>

Mostly research  
(some traction in embedded)  
[RISC-V Instruction Set](#)

# Architecture Sits at the Hardware Interface

## Source code

Different applications or algorithms

## Compiler

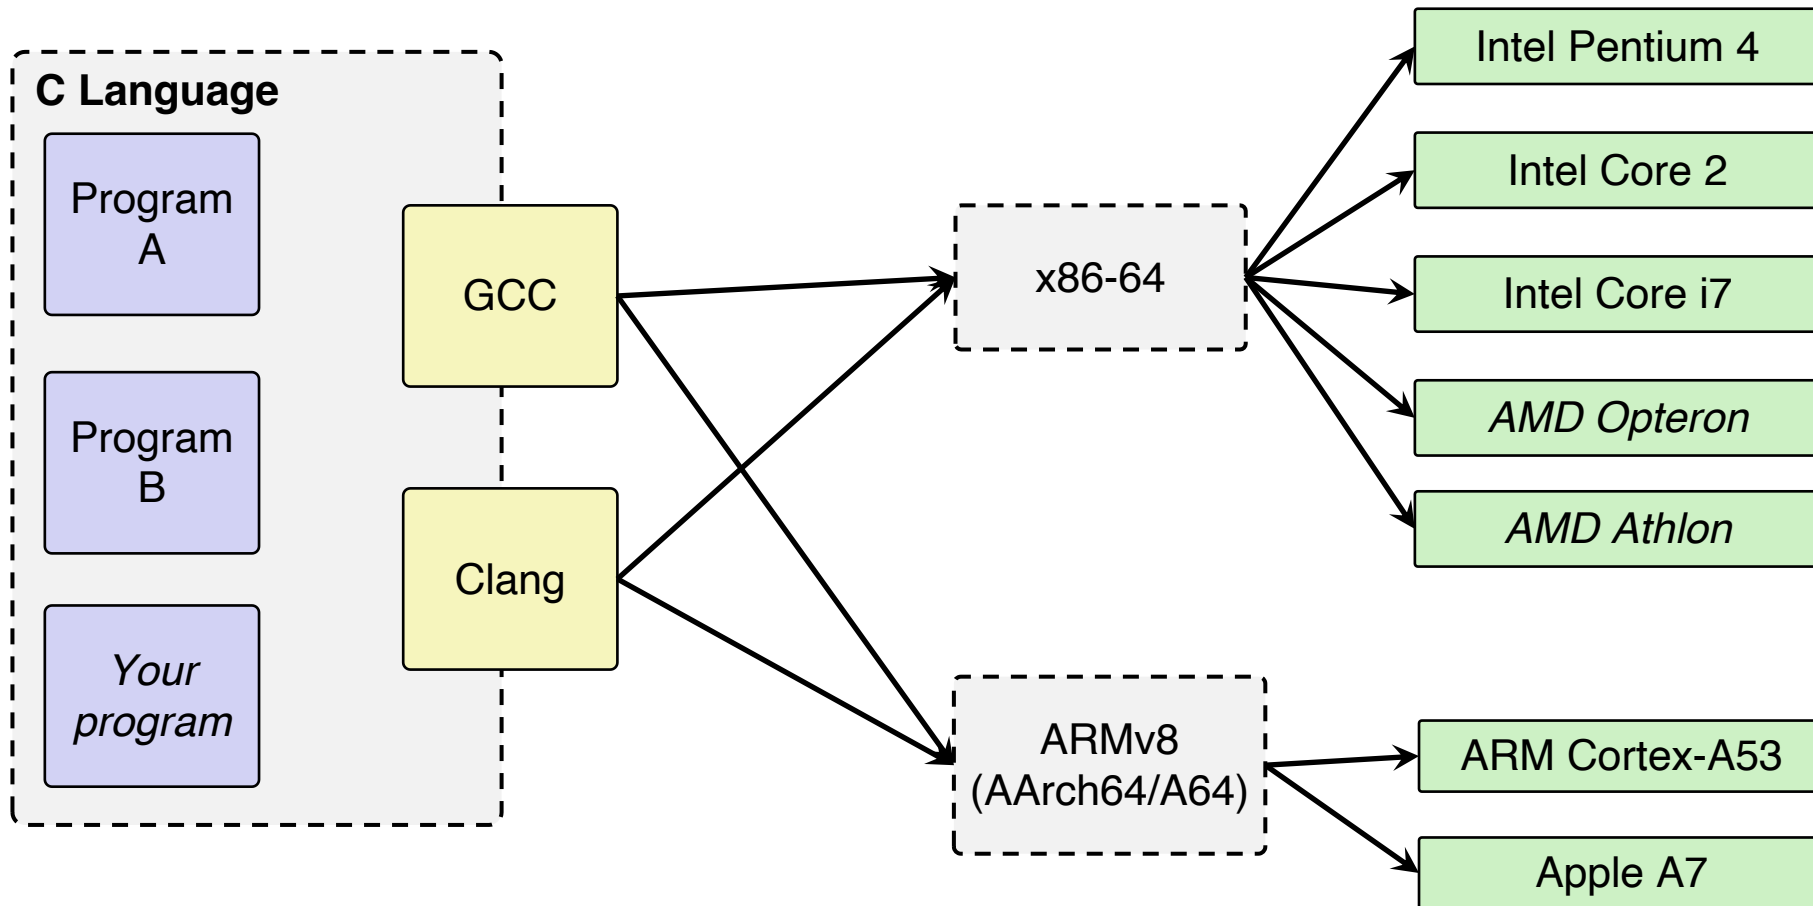
Perform optimizations, generate instructions

## Architecture

Instruction set

## Hardware

Different implementations

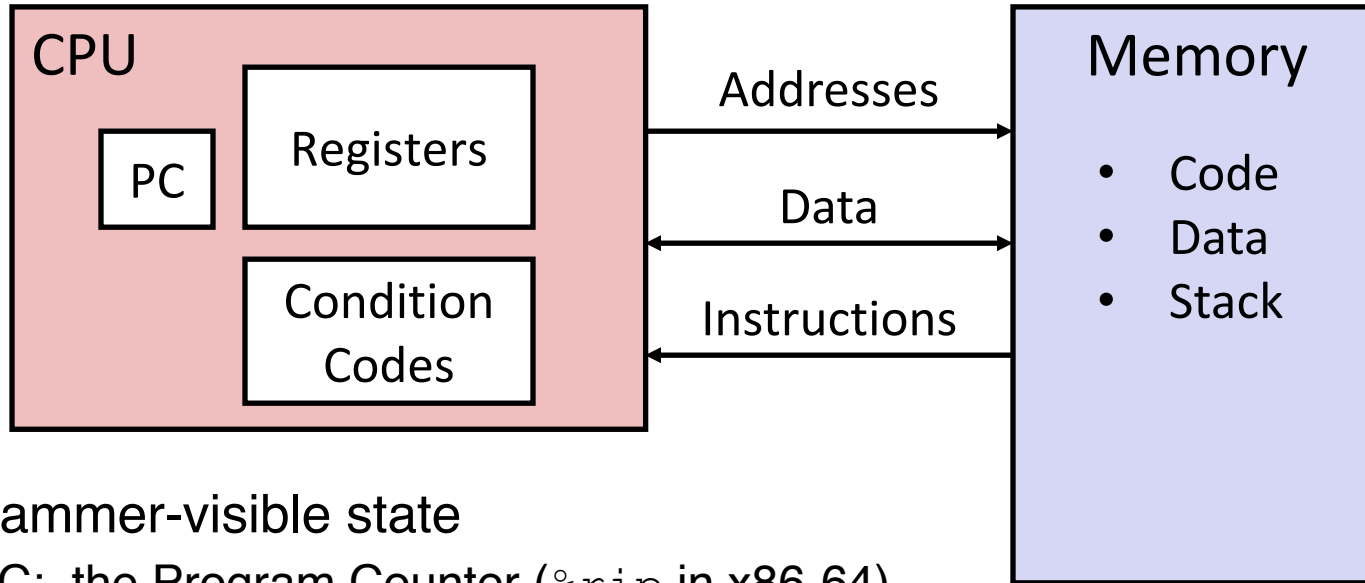


# Writing Assembly Code? In 2024?

- Chances are, you'll never write a program in assembly, but understanding assembly is the key to the machine-level execution model:
  - Behavior of programs in the presence of bugs
    - When high-level language model breaks down
  - Tuning program performance
    - Understand optimizations done/not done by the compiler
    - Understanding sources of program inefficiency
  - Implementing systems software
    - What are the “states” of processes that the OS must manage
    - Using special units (timers, I/O co-processors, etc.) inside processor!
  - Fighting malicious software
    - Distributed software is in binary form



# Assembly Programmer's View



- Programmer-visible state

- PC: the Program Counter (`%rip` in x86-64)
  - Address of next instruction
- Named registers
  - Together in “register file”
  - Heavily used program data
- Condition codes
  - Store status information about most recent arithmetic operation
  - Used for conditional branching

- Memory

- Byte-addressable array
- Code and user data
- Includes *the Stack* (for supporting procedures)

# x86-64 Assembly “Data Types”

- Integral data of 1, 2, 4, or 8 bytes
  - Data values
  - Addresses
- Floating point data of 4, 8, 10 or 2x8 or 4x4 or 8x2
  - Different registers for those (*e.g.*, `%xmm1`, `%ymm2`)
  - Come from *extensions to x86* (SSE, AVX, ...)
- No aggregate types such as arrays or structures
  - Just contiguously allocated bytes in memory
- Two common syntaxes
  - “AT&T”: used by our course, slides, textbook, gnu tools, ...
  - “Intel”: used by Intel documentation, Intel tools, ...
  - Must know which you’re reading

} Not covered  
in 3030

# What is a Register?

- A location in the CPU that stores a small amount of data, which can be accessed very quickly (once every clock cycle)
- Registers have *names*, not *addresses*
  - In assembly, they start with `%` (e.g., `%rsi`)
- Registers are at the heart of assembly programming
  - They are a precious commodity in all architectures, but *especially* x86

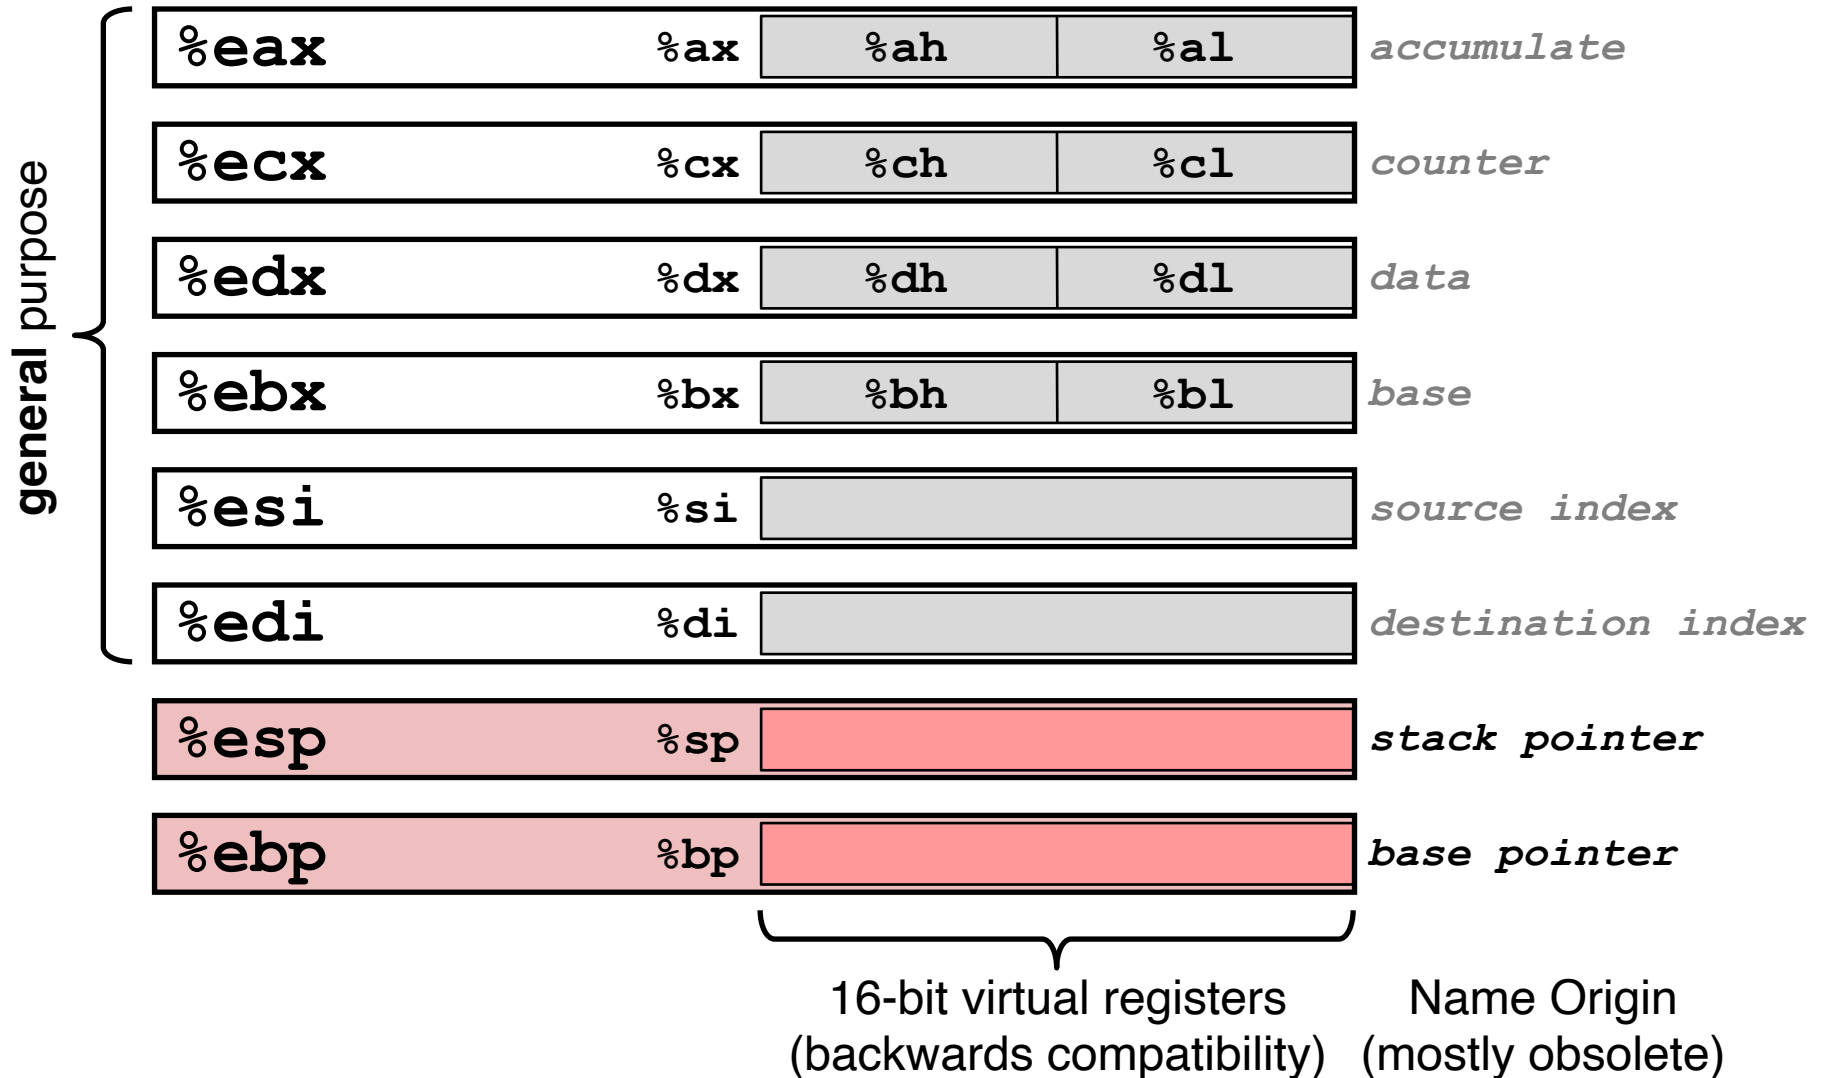
# x86-64 Integer Registers – 64 bits wide

<code>%rax</code>	<code>%eax</code>
<code>%rbx</code>	<code>%ebx</code>
<code>%rcx</code>	<code>%ecx</code>
<code>%rdx</code>	<code>%edx</code>
<code>%rsi</code>	<code>%esi</code>
<code>%rdi</code>	<code>%edi</code>
<code>%rsp</code>	<code>%esp</code>
<code>%rbp</code>	<code>%ebp</code>

<code>%r8</code>	<code>%r8d</code>
<code>%r9</code>	<code>%r9d</code>
<code>%r10</code>	<code>%r10d</code>
<code>%r11</code>	<code>%r11d</code>
<code>%r12</code>	<code>%r12d</code>
<code>%r13</code>	<code>%r13d</code>
<code>%r14</code>	<code>%r14d</code>
<code>%r15</code>	<code>%r15d</code>

- Can reference low-order 4 bytes (also low-order 2 & 1 bytes)

# Some History: IA32 Registers – 32 bits wide



# Memory vs. Register

## Memory

- Addresses
  - `0x7FFFD024C3DC`
- Big
  - $\sim 8$  GiB
- Slow
  - $\sim 50$ - $100$  ns
- Dynamic
  - Can “grow” as needed while program runs

## Register

- vs.** Names
  - `%rdi`
- vs.** Small
  - $(16 \times 8 \text{ B}) = 128 \text{ B}$
- vs.** Fast
  - sub-nanosecond timescale
- vs.** Static
  - fixed number in hardware

# Three Basic Kinds of Instructions

## 1) Transfer data between memory and register

- **Load** data from memory into register
  - $\%reg = \text{Mem}[\text{address}]$
- **Store** register data into memory
  - $\text{Mem}[\text{address}] = \%reg$

**Remember:** Memory is indexed just like an array of bytes!

## 2) Perform arithmetic operation on register or memory data

- $c = a + b;$        $z = x \ll y;$        $i = h \ \& \ g;$

## 3) Control flow: what instruction to execute next

- Unconditional jumps to/from procedures
- Conditional branches

# Instruction Sizes

- Size specifiers

- **b** = 1-byte “byte”  
**l** = 4-byte “long word”

**w** = 2-byte “word”  
**q** = 8-byte “quad word”

- Note that due to backwards-compatible support for 8086 programs (16-bit machines!), “word” means 16 bits = 2 bytes in x86 instruction names



# Operand Types

- **Immediate:** Constant integer data (\$)
  - Examples: \$0x400, \$-533
  - Like C literal, but prefixed with '\$'
  - Encoded with 1, 2, 4, or 8 bytes *depending on the instruction*
- **Register:** 1 of 16 integer registers (%)
  - Examples: %rax, %r13
  - But %rsp reserved for special use
  - Others have special uses for particular instructions
- **Memory:** Consecutive bytes of memory at a computed address (())
  - Simplest example: (%rax)
  - Various other “address modes”

%rax

%rcx

%rdx

%rbx

%rsi

%rdi

%rsp

%rbp

%rN

# x86-64 Introduction

- Data transfer instruction (`mov`)
- Arithmetic operations
- Memory addressing modes
  - `swap` example

# Moving Data

- General form: `mov_ source, destination`
  - Missing letter (`_`) specifies size of operands
  - Lots of these in typical code
- `movb src, dst`
  - Move 1-byte “**byte**”
- `movw src, dst`
  - Move 2-byte “**word**”
- `movl src, dst`
  - Move 4-byte “long word”
- `movq src, dst`
  - Move 8-byte “quad word”

# Operand Combinations

	Source	Dest	Src, Dest	C Analog
movq	Imm	Reg	movq \$0x4, %rax	var_a = 0x4;
		Mem	movq \$-147, (%rax)	*p_a = -147;
	Reg	Reg	movq %rax, %rdx	var_d = var_a;
		Mem	movq %rax, (%rdx)	*p_d = var_a;
	Mem	Reg	movq (%rax), %rdx	var_d = *p_a;

- *Cannot do memory-memory transfer with a single instruction*
  - How would you do it?

# Some Arithmetic Operations

- Binary (two-operand) Instructions:

- Maximum of one  
memory operand

- Beware argument order!
- No distinction between signed and unsigned
  - Only arithmetic vs. logical shifts

Format	Computation	
<b>addq</b> <i>src</i> , <i>dst</i>	$dst = dst + src$	( <i>dst += src</i> )
<b>subq</b> <i>src</i> , <i>dst</i>	$dst = dst - src$	
<b>imulq</b> <i>src</i> , <i>dst</i>	$dst = dst * src$	signed mult
<b>sarq</b> <i>src</i> , <i>dst</i>	$dst = dst \gg src$	Arithmetic
<b>shrq</b> <i>src</i> , <i>dst</i>	$dst = dst \gg src$	Logical
<b>shlq</b> <i>src</i> , <i>dst</i>	$dst = dst \ll src$	(same as <code>salq</code> )
<b>xorq</b> <i>src</i> , <i>dst</i>	$dst = dst \wedge src$	
<b>andq</b> <i>src</i> , <i>dst</i>	$dst = dst \& src$	
<b>orq</b> <i>src</i> , <i>dst</i>	$dst = dst   src$	

↑ operand size specifier

- How do you implement “`rcx = rax + rbx`”?

# Some Arithmetic Operations

- Unary (one-operand) Instructions:

Format	Computation	
<b>incq</b> <i>dst</i>	$dst = dst + 1$	increment
<b>decq</b> <i>dst</i>	$dst = dst - 1$	decrement
<b>negq</b> <i>dst</i>	$dst = -dst$	negate
<b>notq</b> <i>dst</i>	$dst = \sim dst$	bitwise complement

- See CS:APP3e textbook Section 3.5.5 for more instructions:  
`mulq`, `cqto`, `idivq`, `divq`

# Arithmetic Example

Register	Use(s)
%rdi	1 <sup>st</sup> argument (x)
%rsi	2 <sup>nd</sup> argument (y)
%rax	return value

```
long simple_arith(long x, long y)
{
    long t1 = x + y;
    long t2 = t1 * 3;
    return t2;
}
```

```
y += x;
y *= 3;
long r = y;
return r;
```

```
simple_arith:
    addq    %rdi, %rsi
    imulq   $3, %rsi
    movq    %rsi, %rax
    ret
```

# Example of Basic Addressing Modes

```
void swap(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

```
swap:
    movq    (%rdi), %rax
    movq    (%rsi), %rdx
    movq    %rdx, (%rdi)
    movq    %rax, (%rsi)
    ret
```

Compiler Explorer:

<https://godbolt.org/z/zc4Pcq>



# Summary

- x86-64 is a complex instruction set computing (CISC) architecture
  - There are 3 types of operands in x86-64
    - Immediate, Register, Memory
  - There are 3 types of instructions in x86-64
    - Data transfer, Arithmetic, Control Flow